

System Description and Architecture for phase 1

System Overview

The system is a **Key-Value (KV) Store** implemented as a multi-threaded **HTTP Server** with a caching layer and a persistent database backend. It processes simple HTTP requests (PUT, GET, DELETE) to store, retrieve, and remove key-value pairs.

Key Components

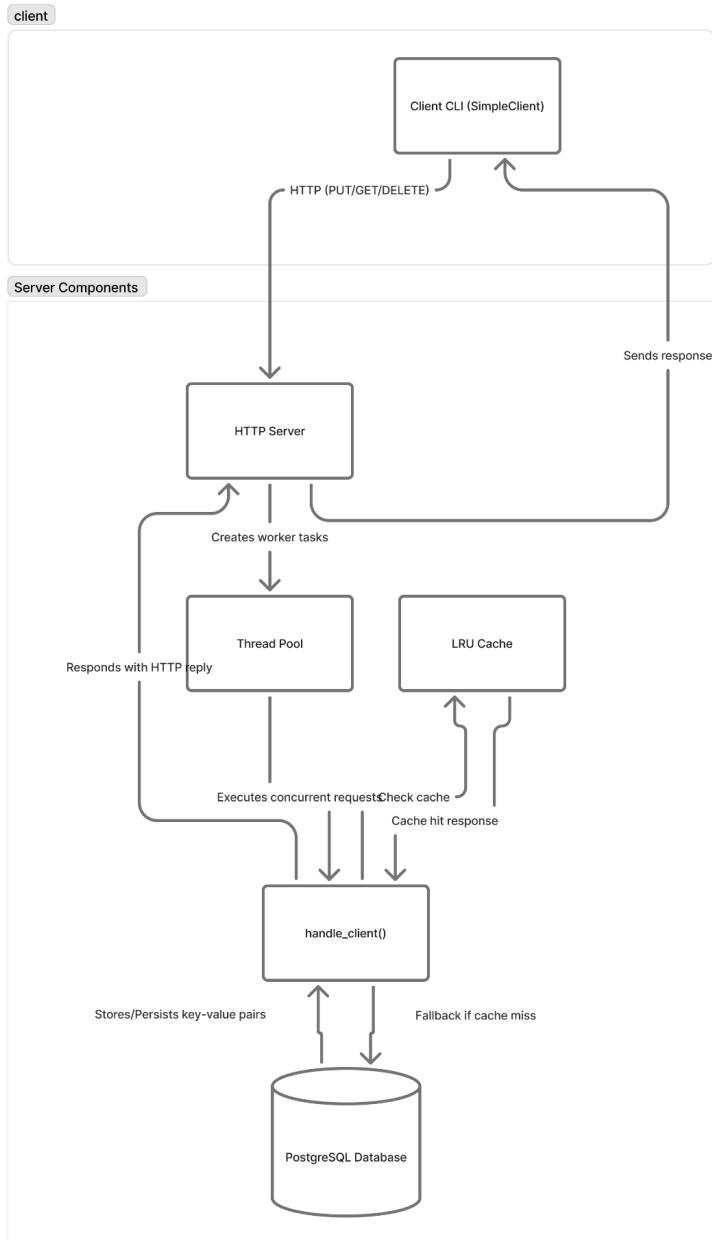
- **SimpleClient (Client Code)**: A basic C++ class using **POSIX sockets** to connect to the server. It crafts and sends raw HTTP/1.1 requests (PUT, GET, DELETE) to the endpoint.
- **HTTPServer (Server Core)**: The main component that sets up the listening socket and manages the application's lifecycle. It uses a **Thread Pool** to handle incoming client connections concurrently.
- **ThreadPool (Concurrency)**: Manages a fixed number of worker threads. It uses a **task queue** and **condition variables** to distribute client handling tasks efficiently among its threads, ensuring high concurrency.
- **LRUCache (In-Memory Cache)**: A thread-safe, memory-backed cache that implements the **Least Recently Used (LRU)** eviction policy. It is used to store frequently accessed data for fast retrieval (cache hits) and reduce database load.
- **Database (Persistence Layer)**: A thread-safe wrapper around the **libpq** library (PostgreSQL client). It handles persistence by storing the key-value pairs in a PostgreSQL database table, ensuring data durability.

Request Flow and Data Management

1. A **Client** sends an HTTP request (e.g., PUT /kv/key, GET /kv/key) to the **HTTPServer**.
2. The **HTTPServer's accept_loop** accepts the connection and **enqueues** the client handling task to the **Thread Pool**.
3. A **Worker Thread** in the pool reads and parses the HTTP request.
 - **PUT Request**: The value is immediately written to both the **LRUCache** and the **Database** (Write-Through strategy).
 - **GET Request**: The system first attempts to retrieve the value from the **LRUCache (Cache Hit)**. If not found (**Cache Miss**), it fetches the value from the **Database**. If the database returns a value (**DB Hit**), the key-value pair is written back to the **LRUCache** (Read-Through strategy) before being returned to the client.
 - **DELETE Request**: The key is removed from both the **LRUCache** and the **Database**.

4. The Worker Thread sends the appropriate HTTP response back to the client.
-

System Architecture Figure



Public GitHub Repository Link

GitHub Link: https://github.com/Pankuu21/kvserver_25m0782